

Reconstruction from Non-Uniform Samples Using a DCT- p Prior

Vinay
SR NO.: 23653

Abstract

This report describes an optimization-based method to restore grayscale images from sparse, noisy pixel measurements. The approach enforces sparsity in the Discrete Cosine Transform domain via a nonconvex ℓ_p -style penalty with $0 < p < 1$. A Majorization-Minimization (MM) routine converts the problem into a sequence of weighted quadratic subproblems. Each subproblem is solved by Preconditioned Conjugate Gradients (PCG). We present the derivation, implementation details, and numerical results on standard test images across multiple sampling ratios and hyperparameters.

1 Introduction

Reconstructing images from incomplete observations is central to many imaging tasks. When only a subset of pixels are available and measurements are noisy, the inverse problem is ill-posed. A common remedy is to impose a prior. Natural images are approximately sparse in transform domains such as the DCT. We exploit this by penalizing DCT coefficients with an ℓ_p -type penalty for $0 < p < 1$. The overall objective balances data fidelity at sampled pixels with a sparsity-promoting regularizer in the DCT domain.

2 Model and algorithm

Let $x^* \in \mathbb{R}^N$ denote the vectorized true image and $m \in \mathbb{R}^M$ the observed sampled pixels. The measurements follow

$$m = Wx^* + \eta,$$

where W extracts sampled pixels and η is additive Gaussian noise. We minimize

$$J(x) = \|Wx - m\|_2^2 + \lambda \sum_{i=1}^N (\varepsilon + y_i^2)^p, \quad y = \text{DCT}(x),$$

with $0 < p < 1$ and small $\varepsilon > 0$.

Because the $(\varepsilon + t)^p$ term is concave in t for $0 < p < 1$, we majorize it using a tangent-line bound at the current iterate. Denote the DCT coefficients at iteration k by $y^{(k)}$. The surrogate becomes a weighted quadratic in x :

$$Q(x | x^{(k)}) = \|Wx - m\|_2^2 + \lambda \sum_i w_i^{(k)} y_i^2 + \text{const},$$

where the weights

$$w_i^{(k)} = p(\varepsilon + (y_i^{(k)})^2)^{p-1}$$

are updated each MM step. Setting $\nabla_x Q = 0$ yields a linear system of the form

$$(W^\top W + \lambda \text{IDCT} \text{diag}(w^{(k)}) \text{DCT})x = W^\top m.$$

We solve this system using PCG with a diagonal (Jacobi) preconditioner for efficiency.

3 Implementation details

The algorithm was implemented in Python using NumPy and SciPy. Key points:

- Images were treated as 256×256 grayscale arrays normalized to $[0, 1]$.
- Random sampling masks were generated for sampling ratios $r \in \{0.1, 0.2, 0.3, 0.5\}$.
- Additive Gaussian noise was scaled to obtain 30 dB SNR on the sampled pixels.
- Parameters swept: $p \in \{0.3, 0.4, 0.5\}$, and $\lambda \in \{10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 1\}$.
- Each MM iteration recomputes weights, assembles the operator via DCT/IDCT and calls PCG.

4 Experimental setup

We ran experiments on two standard test images: *Cameraman* and *Lena*. For each image and sampling ratio:

1. Create a random binary sampling mask with fraction r of pixels observed.
2. Form noisy observations at the sampled pixels with target SNR = 30 dB.
3. Run MM until convergence or until a maximum number of iterations.
4. Record PSNR, ℓ_2 relative error and wall-clock time.

5 Results

5.1 Cameraman

Table 1 summarizes best results (best λ per (r, p)). Visual examples and diagnostics for the best case at $r = 0.5$, $p = 0.5$ are in Figures 1 and 2.

Table 1: Cameraman: best reconstruction metrics for each (r, p) .

r	p	Best λ	Best PSNR (dB)	Rel. Error (ℓ_2)	Time (s)
0.1	0.3	0.1	17.40	0.2548	4.37
0.1	0.4	0.1	18.51	0.2243	3.54
0.1	0.5	0.1	18.57	0.2226	3.39
0.2	0.3	0.1	19.14	0.2084	3.50
0.2	0.4	0.1	20.34	0.1816	3.50
0.2	0.5	0.1	20.38	0.1807	3.44
0.3	0.3	0.1	20.50	0.1783	3.28
0.3	0.4	0.1	21.61	0.1569	3.59
0.3	0.5	0.01	21.83	0.1529	3.45
0.5	0.3	0.01	23.45	0.1269	3.63
0.5	0.4	0.01	24.56	0.1117	3.89
0.5	0.5	0.01	24.88	0.1077	3.36

Visual Results ($r=0.5$, $p=0.5$, $\lambda=0.01$)

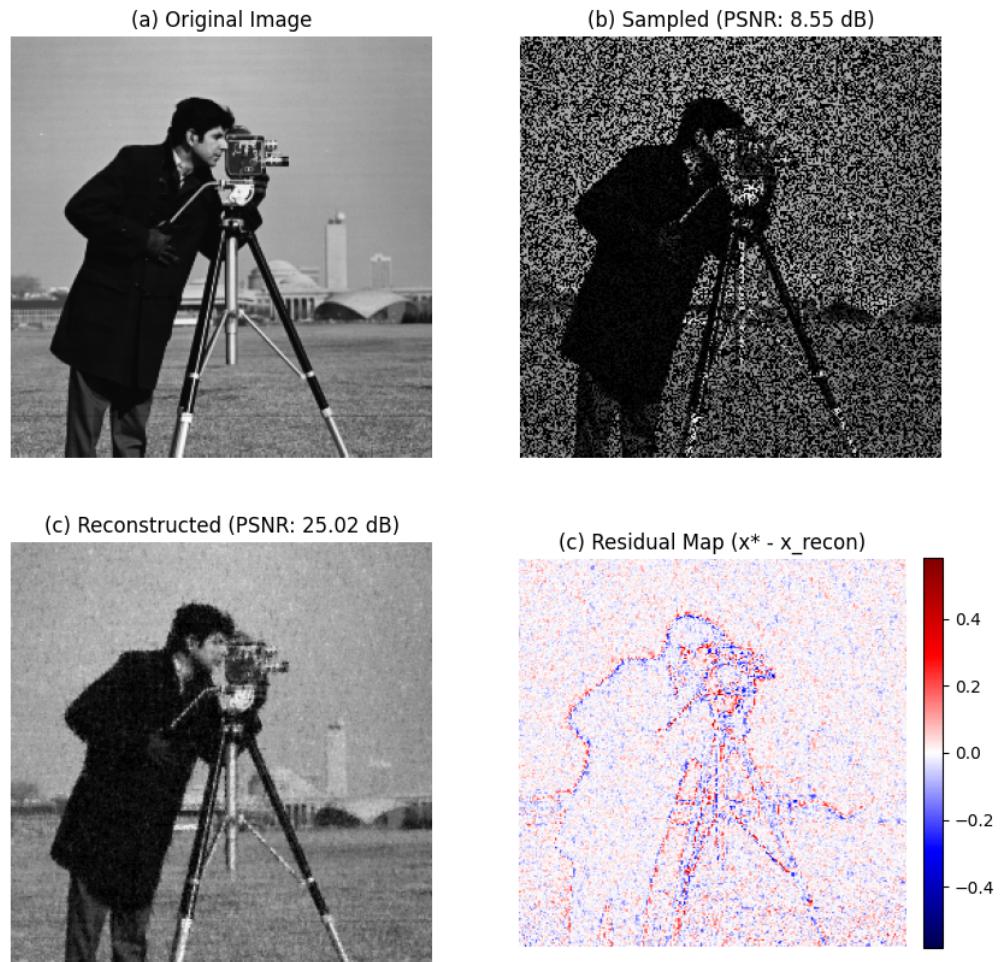


Figure 1: Visual comparison for Cameraman

Diagnostic Plots ($r=0.5$, $p=0.5$, $\lambda=0.01$)

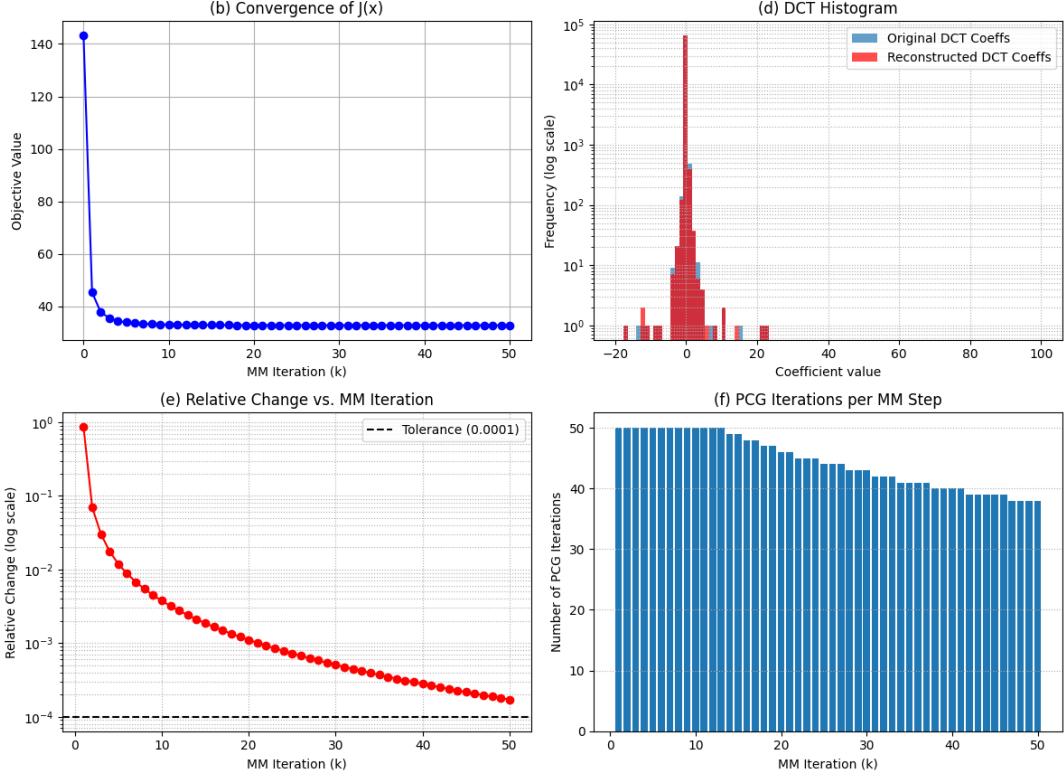


Figure 2: Diagnostic plots for the Cameraman reconstruction at $r=0.5$, $p=0.5$

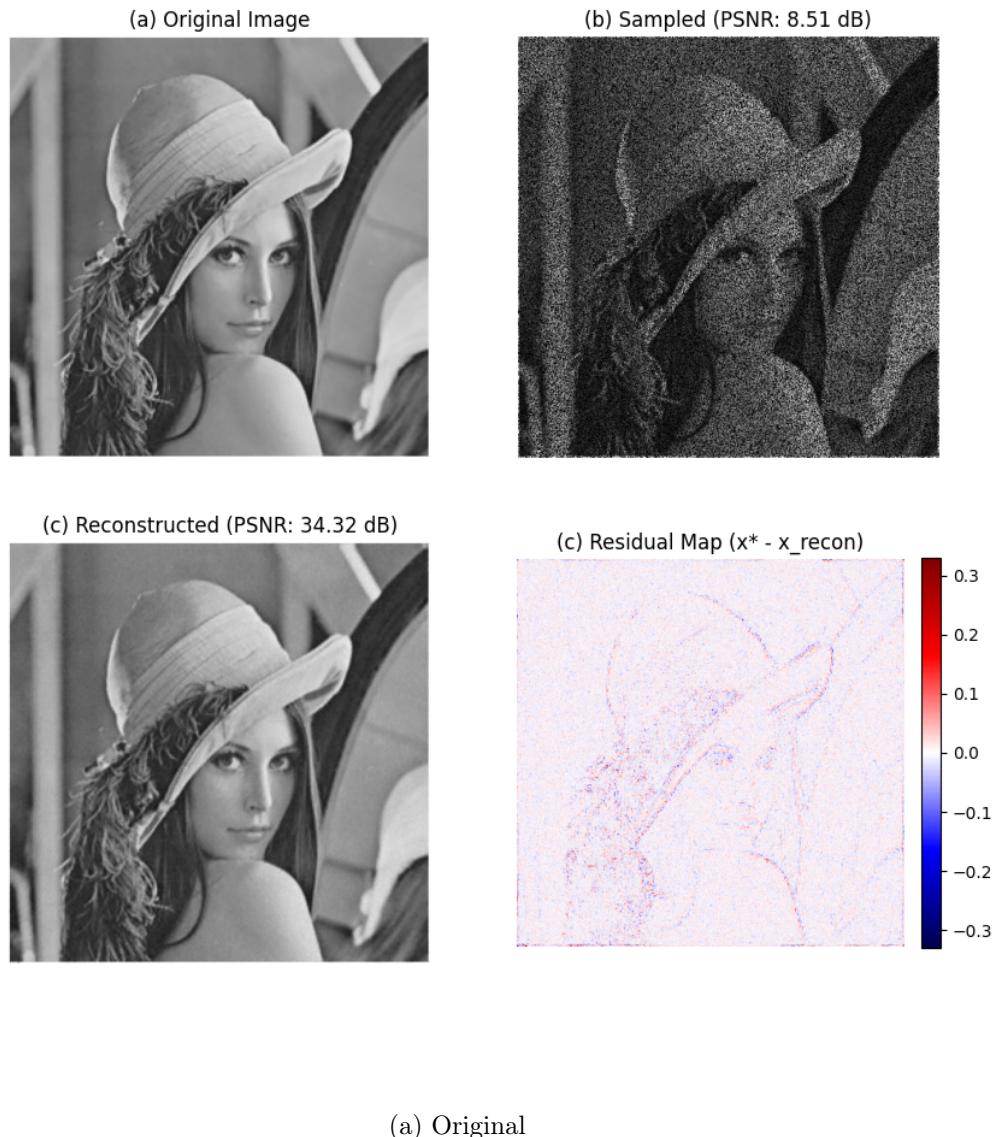
5.2 Lena

Table 2 lists the best-performing λ and metrics for different (r, p) . Representative images and diagnostics for $r = 0.5$, $p = 0.5$ appear in Figures 3 and 4.

Table 2: Lena: best reconstruction metrics for each (r, p) .

r	p	Best λ	Best PSNR (dB)	Rel. Error (ℓ_2)	Time (s)
0.1	0.4	0.1	19.48	0.2262	3.65
0.1	0.5	0.1	19.43	0.2275	3.69
0.2	0.3	0.1	20.54	0.2003	3.40
0.2	0.4	0.1	21.63	0.1766	3.67
0.2	0.5	0.01	22.10	0.1673	4.03
0.3	0.3	0.1	22.10	0.1673	4.03
0.3	0.4	0.01	23.59	0.1409	4.06
0.3	0.5	0.01	23.59	0.1409	4.06
0.5	0.3	0.01	25.93	0.1076	4.26
0.5	0.4	0.01	26.72	0.0983	5.07
0.5	0.5	0.01	27.08	0.0943	3.72

Visual Results ($r=0.5$, $p=0.5$, $\lambda=0.01$)



(a) Original

Figure 3: Lena example (best case at $r = 0.5$, $p = 0.5$).

Diagnostic Plots ($r=0.5$, $p=0.5$, $\lambda=0.01$)

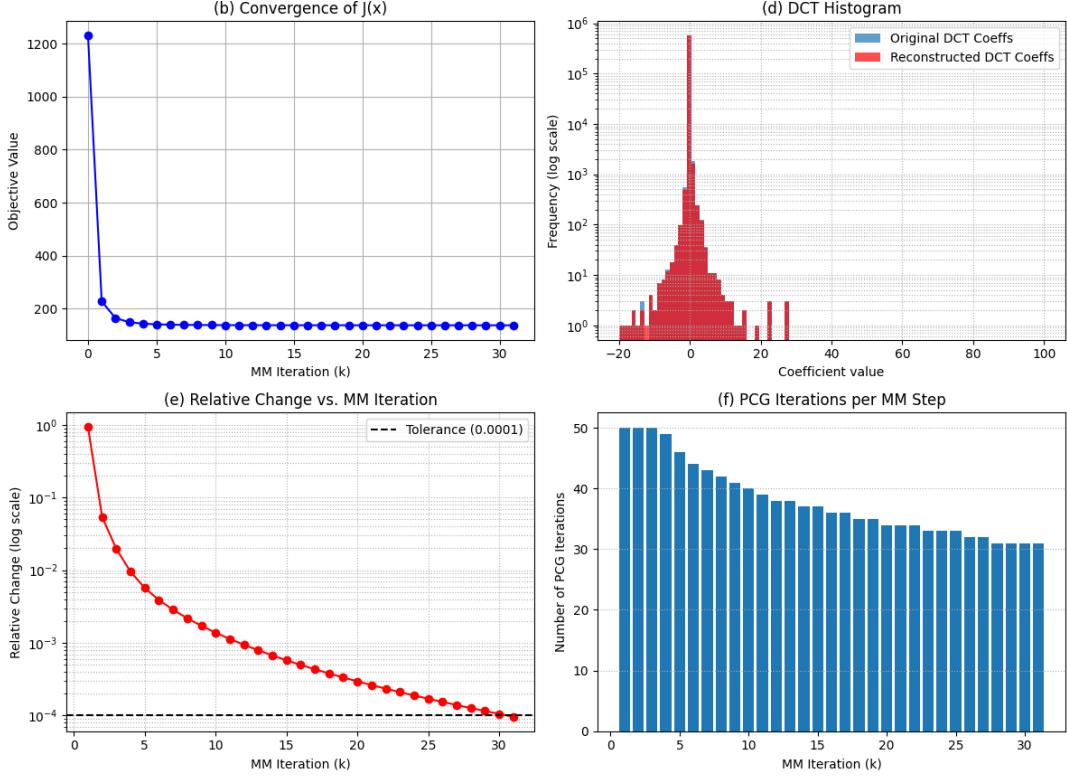


Figure 4: Diagnostic plots for the Lena best case.

6 Discussion

6.1 Sampling fraction

The amount of sampled data is the primary determinant of recovery quality. As the sampling ratio r increases, PSNR improves steadily. This is expected because more direct measurements reduce ambiguity.

6.2 Sparsity exponent p

Larger p values (closer to 1) lead to milder sparsity enforcement. Empirically this often improved PSNR and preserved subtle textures. Very small p values can over-suppress moderate DCT coefficients and harm detail.

6.3 Regularization strength λ

When measured data are scarce, stronger regularization (larger λ) is beneficial. With more samples, a smaller λ lets the algorithm fit observed pixels more closely while relying less on the prior.

6.4 Convergence and artifacts

The MM-PCG scheme converges reliably in a few dozen iterations for the settings used. Sparsity-based priors can introduce small artifacts while filling unobserved regions. These are visible in residual maps and high-frequency DCT bins.

7 Conclusion

We implemented and tested an MM-based solver that enforces DCT-domain sparsity using a nonconvex ℓ_p -style prior. Results on two benchmark images show that reconstruction improves with sampling rate, that moderate p (e.g., 0.5) is often preferable, and that the optimal λ depends on data availability. The approach is computationally efficient when each MM subproblem is solved with PCG and a simple preconditioner.

A Python Implementation Code

This appendix includes the complete Python implementation of the MM-PCG based image reconstruction experiment. All figures and results in the main report were generated using this script.

Source Code

```
1 import numpy as np
2 from PIL import Image
3 from numpy.linalg import norm
4 from scipy.fft import dct, idct
5 import matplotlib.pyplot as plt
6 import time
7 import sys
8
9 # =====
10 # SECTION 1: CORE ALGORITHM FUNCTIONS
11 # =====
12
13 def calculate_psnr(img_true, img_recon):
14     """Calculates the Peak Signal-to-Noise Ratio (PSNR)."""
15     N = img_true.size
16     mse = np.sum((img_true - img_recon)**2) / N
17     if mse == 0:
18         return float('inf')
19     max_intensity = np.max(img_true)
20     psnr_val = 20 * np.log10(max_intensity / np.sqrt(mse))
21     return psnr_val
22
23 def calculate_rel_error(img_true, img_recon):
24     """Calculates the l2 relative error."""
25     return norm(img_recon - img_true) / norm(img_true)
26
27 def calculate_objective(x, m, mask, lam, p_val, epsilon):
28     """Calculates the objective function J(x)."""
29     data_cost = np.sum((mask * x - m)**2)
30     dct_x = dct(dct(x, axis=0, norm='ortho'), axis=1, norm='ortho')
31     sparsity_cost = lam * np.sum((epsilon + dct_x**2)**p_val)
32     return data_cost + sparsity_cost
33
34 def create_random_mask(height, width, sampling_ratio):
35     """Creates a random binary mask."""
36     num_pixels = height * width
37     num_samples = int(num_pixels * sampling_ratio)
38     flat_indices = np.arange(num_pixels)
39     np.random.shuffle(flat_indices)
40     sample_indices = flat_indices[:num_samples]
```

```

41     mask_flat = np.zeros(num_pixels)
42     mask_flat[sample_indices] = 1.0
43     return mask_flat.reshape((height, width))
44
45 def add_snr_noise(image_sampled, mask, target_snr_db=30):
46     """Adds Gaussian noise to achieve a specific SNR."""
47     signal_pixels = image_sampled[mask > 0]
48     signal_norm = np.linalg.norm(signal_pixels)
49     noise_norm = signal_norm / (10**((target_snr_db / 20.0)))
50     noise = np.random.randn(len(signal_pixels))
51     scaled_noise = noise * (noise_norm / np.linalg.norm(noise))
52     m = image_sampled.copy()
53     m[mask > 0] += scaled_noise
54     return m
55
56 def apply_M_operator(z, mask, weights, lam):
57     """Applies the linear operator M(k) to an image z."""
58     part1 = mask * z
59     dct_z = dct(dct(z, axis=0, norm='ortho'), axis=1, norm='ortho')
60     weighted_dct = weights * dct_z
61     part2 = idct(idct(weighted_dct, axis=0, norm='ortho'), axis=1, norm='ortho')
62     return part1 + (lam * part2)
63
64 def preconditioned_cg(b, operator_func, x0, preconditioner, tol=1e-6, max_iter=50):
65     """Solves M(x) = b using Preconditioned Conjugate Gradient."""
66     x = x0.copy()
67     r = b - operator_func(x)
68     z = r / preconditioner
69     p = z.copy()
70     rz_old = np.sum(r * z)
71
72     for i in range(max_iter):
73         Ap = operator_func(p)
74         alpha = rz_old / np.sum(p * Ap)
75         x = x + alpha * p
76         r = r - alpha * Ap
77
78         if np.sqrt(np.sum(r*r)) < tol:
79             break
80
81         z = r / preconditioner
82         rz_new = np.sum(r * z)
83         beta = rz_new / rz_old
84         p = z + beta * p
85         rz_old = rz_new
86
87     return x, i + 1
88
89 def run_reconstruction(m, mask, x_true, lam, p_val, epsilon, max_mm_iter=50,
90     ↪ mm_tol=1e-4, cg_tol=1e-6, max_cg_iter=50):
91     """Runs the full MM-PCG reconstruction and returns history."""
92     x_k = m.copy() #  $x^{(0)}$ 
93
94     history = {
95         'obj': [], 'psnr': [], 'rel_change': [], 'cg_iters': []
96     }
97
98     # Record initial state

```

```

98     obj_val = calculate_objective(x_k, m, mask, lam, p_val, epsilon)
99     psnr_val = calculate_psnr(x_true, x_k)
100    history['obj'].append(obj_val)
101    history['psnr'].append(psnr_val)
102
103    for k in range(max_mm_iter):
104        dct_x = dct(dct(x_k, axis=0, norm='ortho'), axis=1, norm='ortho')
105        weights = p_val * (epsilon + dct_x**2)**(p_val - 1)
106
107        b = m
108        operator_for_cg = lambda z: apply_M_operator(z, mask, weights, lam)
109
110        max_weight = np.max(weights)
111        preconditioner = mask + (lam * max_weight)
112        preconditioner[preconditioner == 0] = 1e-6
113
114        x_k_plus_1, cg_iters = preconditioned_cg(
115            b, operator_for_cg, x_k, preconditioner,
116            tol=cg_tol, max_iter=max_cg_iter
117        )
118
119        relative_change = norm(x_k_plus_1 - x_k) / norm(x_k)
120
121        # Calculate and record diagnostics
122        obj_val = calculate_objective(x_k_plus_1, m, mask, lam, p_val, epsilon)
123        psnr_val = calculate_psnr(x_true, x_k_plus_1)
124        history['obj'].append(obj_val)
125        history['psnr'].append(psnr_val)
126        history['rel_change'].append(relative_change)
127        history['cg_iters'].append(cg_iters)
128
129        x_k = x_k_plus_1
130
131        if relative_change < mm_tol and k > 0:
132            break
133
134        # Add final metrics for easy access
135        history['final_psnr'] = history['psnr'][-1]
136        history['final_rel_err'] = calculate_rel_error(x_true, x_k)
137        history['final_recon'] = x_k
138
139    return history
140
141 def plot_diagnostic_graphs(history, x_true, m, p, r, lam):
142     """
143     Generates the 2x2 diagnostic plots (b, c, d, e, f)
144     and the visual result plots.
145     """
146
147     print("\nGenerating diagnostic plots for best case...")
148
149     x_recon = history['final_recon']
150     psnr_input = calculate_psnr(x_true, m)
151     psnr_final = history['final_psnr']
152     mm_tol = 1e-4 # Define for plot
153
154     # Plot 1: Visual Results (Original, Sampled, Recon, Residual)
155     fig, axes = plt.subplots(2, 2, figsize=(10, 10))

```

```

156     fig.suptitle(f'Visual Results (r={r}, p={p}, $\lambda$={lam})', fontsize=16)
157
158     axes[0, 0].imshow(x_true, cmap='gray', vmin=0, vmax=1)
159     axes[0, 0].set_title(f'(a) Original Image')
160     axes[0, 0].axis('off')
161
162     axes[0, 1].imshow(m, cmap='gray', vmin=0, vmax=1)
163     axes[0, 1].set_title(f'(b) Sampled (PSNR: {psnr_input:.2f} dB)')
164     axes[0, 1].axis('off')
165
166     axes[1, 0].imshow(x_recon, cmap='gray', vmin=0, vmax=1)
167     axes[1, 0].set_title(f'(c) Reconstructed (PSNR: {psnr_final:.2f} dB)')
168     axes[1, 0].axis('off')
169
170     residual = x_true - x_recon
171     vmax = np.max(np.abs(residual))
172     im = axes[1, 1].imshow(residual, cmap='seismic', vmin=-vmax, vmax=vmax)
173     axes[1, 1].set_title(f'(c) Residual Map (x* - x_recon)')
174     axes[1, 1].axis('off')
175     fig.colorbar(im, ax=axes[1, 1], fraction=0.046, pad=0.04)
176     plt.savefig('visuals.png')
177
178 # Plot 2: Diagnostic Plots
179 fig, axes = plt.subplots(2, 2, figsize=(12, 10))
180 fig.suptitle(f'Diagnostic Plots (r={r}, p={p}, $\lambda$={lam})', fontsize=16)
181 mm_iters = np.arange(len(history['obj']))
182
183 # Plot (b) Convergence of J(x)
184 axes[0, 0].plot(mm_iters, history['obj'], 'bo-')
185 axes[0, 0].set_title(f'(b) Convergence of J(x)')
186 axes[0, 0].set_xlabel('MM Iteration (k)')
187 axes[0, 0].set_ylabel('Objective Value')
188 axes[0, 0].grid(True)
189
190 # Plot (d) DCT Histogram
191 dct_true_flat = dct(dct(x_true, axis=0, norm='ortho'), axis=1,
192     ↪ norm='ortho').ravel()
192 dct_recon_flat = dct(dct(x_recon, axis=0, norm='ortho'), axis=1,
193     ↪ norm='ortho').ravel()
193 axes[0, 1].hist(dct_true_flat, bins=100, range=(-20, 100), log=True,
194     ↪ alpha=0.7, label='Original DCT Coeffs')
195 axes[0, 1].hist(dct_recon_flat, bins=100, range=(-20, 100), log=True,
196     ↪ alpha=0.7, label='Reconstructed DCT Coeffs', color='red')
197 axes[0, 1].set_title('(d) DCT Histogram')
198 axes[0, 1].set_xlabel('Coefficient value')
199 axes[0, 1].set_ylabel('Frequency (log scale)')
200 axes[0, 1].legend()
201 axes[0, 1].grid(True, which='both', linestyle=':')
202
203 # Plot (e) Relative Change
204 axes[1, 0].plot(mm_iters[1:], history['rel_change'], 'ro-')
205 axes[1, 0].set_title('(e) Relative Change vs. MM Iteration')
206 axes[1, 0].set_xlabel('MM Iteration (k)')
207 axes[1, 0].set_ylabel('Relative Change (log scale)')
208 axes[1, 0].set_yscale('log')
209 axes[1, 0].axhline(y=mm_tol, color='k', linestyle='--', label=f'Tolerance
210     ↪ ({mm_tol})')
210 axes[1, 0].legend()

```

```

211     axes[1, 0].grid(True, which='both', linestyle=':')
212
213     # Plot (f) PCG Iterations
214     axes[1, 1].bar(mm_iters[1:], history['cg_iters'])
215     axes[1, 1].set_title('(f) PCG Iterations per MM Step')
216     axes[1, 1].set_xlabel('MM Iteration (k)')
217     axes[1, 1].set_ylabel('Number of PCG Iterations')
218     axes[1, 1].grid(True, axis='y', linestyle=':')
219
220     plt.tight_layout(rect=[0, 0.03, 1, 0.95])
221     plt.savefig('diagnostic_plots.png')
222
223 def save_table_as_figure(data_list_of_dicts, col_headers, title, filename):
224     """
225     Saves the final results table as a PNG image.
226     """
227     print(f"\nGenerating table image: {filename}...")
228
229     # 1. Convert data from list-of-dicts to list-of-lists (strings)
230     cell_text = []
231     for run_data in data_list_of_dicts:
232         row = [
233             f"{run_data['r']:.1f}",
234             f"{run_data['p']:.1f}",
235             f"{run_data['best_lambda']:.2g}", # Use .2g for 0.01, 0.1
236             f"{run_data['best_psnr']:.2f}",
237             f"{run_data['rel_err']:.4f}",
238             f"{run_data['runtime']:.2f}"
239         ]
240         cell_text.append(row)
241
242     # 2. Create the figure and table
243     # Adjust figsize; (width, height)
244     fig, ax = plt.subplots(figsize=(12, len(cell_text) * 0.4 + 1))
245     ax.axis('off') # Hide axes (x, y)
246     ax.axis('tight')
247
248     # Create the table
249     table = ax.table(cellText=cell_text,
250                       colLabels=col_headers,
251                       loc='center',
252                       cellLoc='center')
253
254     # 3. Style the table
255     table.auto_set_font_size(False)
256     table.set_fontsize(10)
257     table.scale(1.1, 1.4) # Adjust scale (width, height)
258
259     # Style header row
260     for (i, j), cell in table.get_celld().items():
261         if i == 0: # Header row
262             cell.set_text_props(weight='bold')
263
264     # 4. Add title
265     plt.title(title, weight='bold', fontsize=12, y=1.08)
266
267     # 5. Save the figure
268     plt.savefig(filename, bbox_inches='tight', dpi=200, pad_inches=0.1)

```

```

269     plt.close(fig) # Close the figure to free memory
270     print(f"Saved table to {filename}")
271
272 # =====
273 # SECTION 2: MAIN EXPERIMENT SCRIPT
274 # =====
275 if __name__ == '__main__':
276
277     # --- 1. Define All Experiment Parameters ---
278
279     IMAGE_PATH = 'images/lena.png'
280     EPSILON = 1e-6
281     MM_TOL = 1e-4
282     MAX_MM_ITER = 50 # Set a fixed number for consistent timing
283
284     # Parameters to sweep
285     r_values = [0.1, 0.2, 0.3, 0.5]
286     p_values = [0.3, 0.4, 0.5]
287     lambda_values = [1e-4, 1e-3, 1e-2, 1e-1, 1.0]
288
289     col_headers = ["Sampling (r)", "p-value", "Best Lambda",
290                    "Best PSNR (dB)", "Rel. Error (12)", "Runtime (s)"]
291
292     # --- 2. Load the Original Image ---
293     try:
294         x_true = np.array(Image.open(IMAGE_PATH).convert('L')) / 255.0
295     except FileNotFoundError:
296         print(f"Error: Image '{IMAGE_PATH}' not found.")
297         sys.exit()
298
299     print(f"--- Running Full Experiment Sweep for '{IMAGE_PATH}' ---")
300
301     # Print the table header
302     print("*" * 70)
303     print(f"{'Sampling (r)':<12} | {'p-value':<7} | {'Best Lambda':<11} | {'Best PSNR'>
304          (dB)':<15} | {'Rel. Error (12)':<15} | {'Runtime (s)':<10}")
305     print("-" * 70)
306
307     # --- 3. Start the Triple Loop ---
308
309     all_best_results = []
310     # This will store the data needed for Plot (a)
311     psnr_vs_lambda_data_r0_5_p0_5 = []
312     # This will store the full history for the best run for Plots (b-f)
313     best_history_r0_5_p0_5 = None
314     best_lambda_for_plots = None
315     best_psnr_for_plots = -1.0
316
317     # --- Loop 1: Sampling Ratio (r) ---
318     for r in r_values:
319
320         # Create mask and noisy data ONCE for this r
321         mask = create_random_mask(x_true.shape[0], x_true.shape[1], r)
322         x_sampled = x_true * mask
323         m = add_snr_noise(x_sampled, mask, target_snr_db=30)
324
325         # --- Loop 2: p-value ---
326         for p in p_values:

```

```

326
327     sweep_results = []
328
329     # --- Loop 3: Lambda ( $\lambda$ ) ---
330     for lam in lambda_values:
331
332         print(f"  Running: r={r}, p={p}, \u03bb={lam:.1e}...")
333
334         start_time = time.time()
335
336         history = run_reconstruction(
337             m, mask, x_true,
338             lam=lam, p_val=p, epsilon=EPSILON,
339             max_mm_iter=MAX_MM_ITER, mm_tol=MM_TOL
340         )
341
342         end_time = time.time()
343         runtime = end_time - start_time
344
345         # Store the results of this single run
346         run_data = {
347             'lambda': lam,
348             'psnr': history['final_psnr'],
349             'rel_err': history['final_rel_err'],
350             'runtime': runtime,
351             'history': history # Store the full history
352         }
353         sweep_results.append(run_data)
354
355         # --- Special step: Save data for the plots ---
356         # If this is the specific case we want to plot (r=0.5, p=0.5)
357         if r == 0.5 and p == 0.5:
358             psnr_vs_lambda_data_r0_5_p0_5.append(run_data)
359
360             # Check if this is the best PSNR *for this case*
361             if run_data['psnr'] > best_psnr_for_plots:
362                 best_psnr_for_plots = run_data['psnr']
363                 best_history_r0_5_p0_5 = history
364                 best_lambda_for_plots = lam
365
366             # --- Find the BEST result from the lambda sweep ---
367             best_run = max(sweep_results, key=lambda x: x['psnr'])
368             table_row_data = {
369                 'r': r,
370                 'p': p,
371                 'best_lambda': best_run['lambda'],
372                 'best_psnr': best_run['psnr'],
373                 'rel_err': best_run['rel_err'],
374                 'runtime': best_run['runtime']
375             }
376             # Store the best one for the table
377             all_best_results.append(table_row_data)
378
379             # Print the row for the table
380             print(f"{r:<12.1f} | {p:<7.1f} | {best_run['lambda']:<7.2g} |"
381                 f" {best_run['psnr']:<15.2f} | {best_run['rel_err']:<15.4f} |"
382                 f" {best_run['runtime']:<10.2f}")

```

```

382     print("=="*80)
383     print("--- Experiment Sweep Complete ---")
384
385     # --- 4. Generate Plot (a) ---
386     print("\nGenerating 'psnr_vs_lambda_plot.png'...")
387
388     # Extract data for plotting
389     lambdas = [run['lambda'] for run in psnr_vs_lambda_data_r0_5_p0_5]
390     psnrs = [run['psnr'] for run in psnr_vs_lambda_data_r0_5_p0_5]
391
392     plt.figure(figsize=(8, 6))
393     plt.plot(lambdas, psnrs, 'bo-')
394     plt.title(f'(a) PSNR vs. $\lambda$ (r=0.5, p=0.5)')
395     plt.xlabel('$\lambda$ (Lambda)')
396     plt.ylabel('PSNR (dB)')
397     plt.xscale('log')
398     plt.grid(True, which='both', linestyle=':')
399     plt.savefig('psnr_vs_lambda_plot.png')
400
401     # --- 5. Generate Plots (b, c, d, e, f) ---
402     if best_history_r0_5_p0_5:
403         # Re-create the mask and noise for this specific case to plot
404         r_plot, p_plot = 0.5, 0.5
405         mask_plot = create_random_mask(x_true.shape[0], x_true.shape[1], r_plot)
406         x_sampled_plot = x_true * mask_plot
407         m_plot = add_snr_noise(x_sampled_plot, mask_plot, target_snr_db=30)
408
409         plot_diagnostic_graphs(best_history_r0_5_p0_5,
410                                x_true, m_plot,
411                                p=p_plot, r=r_plot, lam=best_lambda_for_plots)
412         print("Generated 'final_reconstruction_visuals.png' and"
413               " 'final_diagnostic_plots.png'")
414
415     # --- 6. NEW: Save the final table as a PNG ---
416     table_title = f'Table: Best Reconstruction Metrics for {IMAGE_PATH}'
417     save_table_as_figure(all_best_results, col_headers, table_title,
418                           'results_table.png')
419
420     print("\nAll tasks complete.")

```

Repository Link

The complete implementation, figures, and dataset are available on GitHub:

<https://github.com/vinaaaaay/LN0-Project.git>